

# RPC: Remote Procedure Call

February 24, 2017

# Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

Further Reading

# Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults


Further Reading

# Remote Procedure Call (RPC)


- ▶ Message-based programming with `send()`/`receive()` primitives is not convenient
  - ▶ depends on the communication protocol used (TCP vs. UDP)
  - ▶ requires the specification of an application protocol
  - ▶ akin to I/O
- ▶ Function/procedure call in a remote computer
  - ▶ is a familiar paradigm
  - ▶ eases transparency
  - ▶ is particularly suited for client-server applications

# RPC: the Idea

## Local procedure call:




## Remote procedure call:




# Program Development with RPCs: the Vision

- Design/develop an application ignoring distribution



- Distribute *a posteriori*



# Roadmap

Idea

**Implementation**

Transparency

RPC Semantics in the Presence of Faults

Further Reading


# RPC Stub Routines

- ▶ Ensure RPC transparency

**Client** invokes the **client stub** – a local function


**Remote function** is invoked by the **server stub** – a local function

- ▶ The stub routines communicate with one another by exchanging messages






# Well Known Trick: also Used for System Calls



# Typical Architecture of an RPC System



**Obs.** RPC is typically implemented on top of the transport layer (TCP/IP)

# Client Stub

## Request

1. Assembles message: **parameter marshalling**
2. Sends message, via `write()` / `sendto()` to server
3. Blocks waiting for response, via `read()` / `recvfrom()`
  - ▶ Not in the case of **asynchronous RPC**

## Response

1. Receives responses
2. Extracts the results (**unmarshalling**)
3. Returns to client
  - ▶ Assuming **synchronous RPC**

# Server Stub

## Request


1. Receives message with request, via `read()` / `recvfrom()`
2. Parses message to determine arguments (**unmarshalling**)
3. Calls function

## Response

1. Assembles message with the return value of the function
2. Sends message, via `write()` / `sendto()`
3. Blocks waiting for a new request

# RPC: Dispatching

- ▶ Often, **RPC services** offer more than one remote procedure:



- ▶ The identification of the procedure is performed by the **dispatcher**
  - ▶ This leads to a hierarchical name space (**service, procedure**)

# Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

Further Reading

# Transparency: Platform Heterogeneity

Problems at least two:

1. Different architectures use different formats
  - ▶ 1's-complement vs. 2's complement
  - ▶ big-endian vs. little-endian
  - ▶ ASCII vs. UTF-??
2. Languages or compilers may use different representations for composite data-structures

Solution mainly two:

standardize format in the wires

- + needs only two conversions in each platform
- may not be efficient

receiver-makes-right

# Transparency: Addresses as Arguments

**Issue** The meaning of an address (C pointer) is specific to a process

**Solution** Use **call-by-copy/restore** for parameter passing

- + Works in most cases
- Complex
  - ▶ The same address may be passed in different arguments
- Inefficient
  - ▶ For complex data structures, e.g. trees



# Transparency in the Presence of Faults

Problem What if something breaks?

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
  - ▶ Must use request identifiers (sequence nos.)
  - ▶ Must save most recent responses for replay, if the request is not **idempotent**

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
  - ▶ Must use request identifiers (sequence nos.)
  - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes



# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
  - ▶ Must use request identifiers (sequence nos.)
  - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
  - ▶ Was the request processed before the crash?

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
  - ▶ Must use request identifiers (sequence nos.)
  - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
  - ▶ Was the request processed before the crash?
- ▶ Client crashes

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
  - ▶ Must use request identifiers (sequence nos.)
  - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
  - ▶ Was the request processed before the crash?
- ▶ Client crashes
  - ▶ Need to prevent **orphan** computations, i.e. on behalf of a dead process.

# Transparency in the Presence of Faults

**Problem** What if something breaks?

- ▶ The client cannot locate the server
  - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
  - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
  - ▶ Must use request identifiers (sequence nos.)
  - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
  - ▶ Was the request processed before the crash?
- ▶ Client crashes
  - ▶ Need to prevent **orphan** computations, i.e. on behalf of a dead process.

**Issue** A client cannot distinguish between loss of a request, loss of a response or a server crash

- ▶ The absence of a response may be caused by a slow network/server

# Roadmap

Idea

Implementation

Transparency

**RPC Semantics in the Presence of Faults**

Further Reading

# RPC Semantics in the Presence of Faults (Spector82)

**Question** What can a client expect when there is a fault?

**Answer** Depends on the semantics in the presence of faults provided by the RPC system

**At-least-once** Client stub must keep retransmitting until it obtains a response


- ▶ Be careful with non-idempotent operations
- ▶ Spector allows for zero executions in case of server failure

**At-most-once** Not trivial if you use a non-reliable transport, e.g. UDP.

- ▶ If the RPC uses TCP, it may report an error when the TCP connection breaks

**Exactly-once** Not always possible to ensure this semantics, especially if there are external actions that cannot be undone

# Faults and Exactly-once Semantics



**Problem** In the case of external actions, e.g. file printing, it is virtually impossible to ensure Exactly-once Semantics

**Server strategy** One of two:

1. Send an `ACK` after printing
2. Send an `ACK` before printing

**Client strategy** One of four:

1. Never resend the request
2. Always resend the request
3. Resend the request when it receives an `ACK`
4. Resend the request when it does not receive an `ACK`

# Server Faults and Exactly-once Semantics

**Scenario** Server crashes and quickly recovers so that it is able to handle client retransmission, but **it has lost all state**

Let

A: ACK

P: print

C: crash

**Fault scenarios** (ACK $\rightarrow$ P)

1. A $\rightarrow$ P $\rightarrow$ C
2. A $\rightarrow$ C ( $\rightarrow$ P)
3. C ( $\rightarrow$ A $\rightarrow$ P)

**Fault scenarios** (P $\rightarrow$ ACK)

1. P $\rightarrow$ A $\rightarrow$ C
2. P $\rightarrow$ C ( $\rightarrow$ A)
3. C ( $\rightarrow$ P $\rightarrow$ A)

## Client

### Reissue Strategy

Always
Never
When Ack
When not Ack

OK = Text printed once

### Strategy A $\rightarrow$ P

APC	AC(P)	C(AP)
Dup	OK	OK
OK	Zero	Zero
Dup	OK	Zero
OK	Zero	OK

Dup = Text printed twice

## Server

### Strategy P $\rightarrow$ A

PAC	PC(A)	C(PA)
Dup	Dup	OK
OK	OK	Zero
Dup	OK	Zero
OK	Dup	OK

Zero = Text not printed at all

**Conclusion** No combined strategy works on every fault scenario



At-least-once versus at-most-once?

let's take an example: acquiring a lock

if client and server stay up and no message loss, client receives lock

if client fails, it may know it has the lock or not (server needs a plan!)

if server fails, client may have lock or not

at-most-once: client may receive an exception

at-least-once: client keeps trying until, it has the lock or gives up

what does a client do in the case of an exception?

need to implement some application-specific protocol

ask server, do i have the lock?

server needs to have a plan for remembering state across reboots

e.g., store locks on disk.

at-least-once

clients keep trying. server may run procedure several times

server must use application state to handle duplicates

if requests are not idempotent

but difficult to make all request idempotent

e.g., server could store on disk who has lock and req id

check table for each request

even if server fails and reboots, we get correct semantics

Which is right?

depends where RPC is used.

simple applications:

at-least-once is easiest if procedure is idempotent

at-most-once is cool (more like procedure calls)

more sophisticated applications:

need an application-level plan in both cases

not clear at-most-once gives you a leg up

=> Need to handle failures makes RPC very different than procedure calls

src: based F. Kaashoek (according to Dave Andersen)

# Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

Further Reading

# Further Reading

- ▶ Tanenbaum e van Steen, *Distributed Systems, 2nd Ed.*
  - ▶ Section 4.2 *Remote Procedure Call*, except subsection 4.2.4
  - ▶ Subsection 8.3.2 *RPC Semantics in the Presence of Failures*
- ▶ Birrel and Nelson, "*Implementing Remote Procedure Calls*", ACM Transactions on Computer Systems, Vol. 2, No. 1, February 1984, Pages 39-59
- ▶ A. Spector, "*Performing Remote Operations Efficiently on a Local Computer Network*", Communications of the ACM, Vol. 25, No. 4, April 1982, Pages 246-260