

Clients and Servers (Processing)

March 2, 2017

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures


Security

Communication Channel Adaptation


Further Reading

Clientes e Servidores

- ▶ Most distributed applications have a **client-server** architecture:



- ▶ We'll use *client* and *server* in a broad sense:



- ▶ A server can also play the role of client of another service.

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Server/Object Location

Problem: how does a client find a server?

Solution: not one, but several alternatives:

- ▶ hard coded, rarely;
- ▶ program arguments: more flexible, but ...
- ▶ configuration file
- ▶ via *broadcast/multicast*;
- ▶ via location/naming server (later in the course)
 - ▶ local, like `portmapper` or `rmiregistry`;
 - ▶ global.

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading


Distribution Transparency

Issue: Many distribution transparency facets can be achieved through client side **stubs** (also called **clerks**):

Access e.g. via RPC;

Location e.g. via multicast;

Replication e.g. by invoking operations on several replicas:



Faults e.g. by masking server and communication faults

► if possible

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Concurrency


- ▶ There are several reasons for using concurrency:
 - ▶ Performance (+ on servers);
 - ▶ Usability (+ on clients) – still performance, really.
- ▶ The goal is to overlap I/O with processing
- ▶ Example: Web service

Client-side

- ▶ A Web page may be composed of several *objects*
- ▶ A browser can render some objects, while it fetches others via the net.

Server-side

- ▶ May serve several requests simultaneously



src:Pai et al. 99

How to Achieve Concurrency?

Threads

- ▶ Remember SO ...

Events

- ▶ Remember LCOM ...


Iterative Web Server




src:Pai et al. 99

- ▶ Has only one thread
- ▶ Processes a request/connection at a time
- ▶ To read the file the server may have to go to disk. In that case, it:
 - ▶ will block and
 - ▶ cannot process other requests
- ▶ Such a server can process only a few requests per time unit

Multi-threaded Server



- ▶ Each thread processes a request (and HTTP 1.0 connection)
- ▶ If the number of threads is larger than the number of cores/processors
 - ▶ When one thread blocks on I/O
 - ▶ Another thread may be scheduled to run in its place.
- ▶ A common pattern is:
 - One dispatcher thread, which accepts the requests
 - Several worker threads, which process the requests

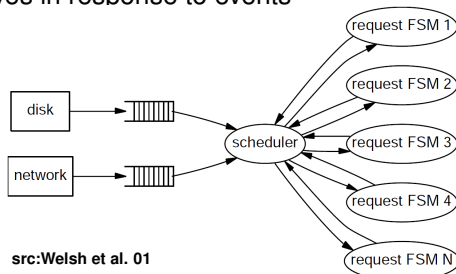


Event-driven Server



src:Pai et al. 99

- ▶ The server executes a loop, in which it:
 - ▶ waits for events (usually I/O events)
 - ▶ processes these events (sequentially)
- ▶ Blocking is avoided by using **non-blocking** I/O operations
- ▶ Known as the state machine approach
 - ▶ The state of the server evolves in response to events
- ▶ For more complex services:
 - ▶ Each request is a FSM
 - ▶ The loop dispatches the event to the appropriate FSM



src:Welsh et al. 01

Thread vs. Event Debate

Ease of programming
Performance

Thread-based Concurrency: Ease of Programming


- ▶ Appears simple:
 - ▶ Structure of each thread similar to that of an iterative server
 - ▶ Need **only** to ensure **isolation** in the access to shared data structures
- ▶ Could use only monitors and condition variables, e.g. synchronized methods in Java
 - ▶ Not so easy: there are some implications in terms of modularity ([Ousterhout96](#))
 - ▶ Possibility of deadlocks
- ▶ Performance may suffer
 - ▶ The larger the critical sections, less concurrency
 - ▶ But the main reason for concurrency is performance

Event-based Concurrency: Ease of Programming

- ▶ Programmer needs to:
 - ▶ Break processing according to potentially blocking calls
 - ▶ Manage the state explicitly (using state machines), rather than relying on the stack
- ▶ The structure of the code is very different from that of the iterative server
- ▶ No nasty errors like race conditions, which may be elusive
- ▶ But many complain about lack of support by debugging tools
- ▶ ... and **others** that the it leads to poorly structured code
 - ▶ Actually, the author interestingly points out the issue is preemption rather than multithreading

Thread-Based Concurrency: Performance

- ▶ Same file 8 KB reads (no disk accesses)
- ▶ No thread creation
- ▶ "4-way 500MHz Pentium III with 2 GB memory under Linux 2.2.14"




src: Welsh et al. 01

- ▶ As the number of threads increases, the system throughput increases, then levels-off and finally dives
- ▶ Clearly each thread requires some resources
- ▶ There are also issues concerning context switching
 - ▶ Actually, depend on whether user-level or kernel-level threads

Event-Based Concurrency: Performance

- ▶ Requires non-blocking (also called asynchronous) I/O operations
 - ▶ Otherwise, must resort to multiple threads to emulate asynchronous I/O
- ▶ Allows user level scheduling
 - ▶ The dispatcher may choose which event to handle next

- ▶ Same file 8 KB reads (no disk accesses)
- ▶ Only one thread
- ▶ As the number of requests in a queue increases throughput increases until it reaches a plateau



src: Welsh et al. 01

- ▶ Needs multiple threads to achieve **parallelism** in multi core/processor platforms

TB vs EB Concurrency: Performance

- ▶ The debate was somewhat "muddled" by implementations that were less than optimal
- ▶ Actually, at the technical level this is very similar to the debate about user-level vs. kernel-level threads
- ▶ User-level threads are more efficient than kernel-level threads
 - ▶ Function calls vs. system calls
 - ▶ But efficient implementations require OS support for non-blocking I/O
- ▶ But there are some unavoidable blocking, e.g. page faults
- ▶ We need kernel-level threads in order to take advantage of multiple processors/cores

Server Architectures

Architecture	Paral.	I/O Oper.	Progr.
Iterative	No	Blocking	easy
Multi-threaded	Yes	Blocking	races
State-machine	Yes	Non-blocking	event-driven

- ▶ To take advantage of multiple processors/cores we need to use *kernel-level threads* (or processes).
 - ▶ On state-machine designs we may use multiple threads

TB vs EB Concurrency: Conclusion

- ▶ Pure thread-based and event-based designs are the extremes in a design space
- ▶ Threads are not as heavy as processes, but they still require resources
 - ▶ You may want to bound their number
- ▶ If you want more parallelism, you need to use an event-based design
- ▶ There are many frameworks for supporting event-driven designs
 - ▶ Java itself offers Java NIO (non-blocking I/O)
- ▶ Not sure about their performance
 - ▶ They are often built on top of a stack of multiple layers

Thread-based Concurrency: Practical Considerations

Java


- ▶ Assume that the Java socket API is not thread-safe
 - ▶ The documentation is mute about this
 - ▶ Java runs on top of different OS
- ▶ You must handle concurrency explicitly

POSIX

- ▶ It requires many system calls, such as `accept`, `read/write`, `sendto/receivefrom`, to be **thread-safe**
 - ▶ But, data of concurrent `write`'s may be interleaved
 - ▶ I.e., `write/read` may not be **atomic** (apparently it depends on the buffer size)
- ▶ What about `send(to)/receive(from)` ?
 - ▶ When used on `STREAM` sockets, may behave similarly to `write`
 - ▶ When used on `DATAGRAM` sockets, one expects POSIX-atomicity to be implied, but ...
- ▶ To be on the safe side, handle concurrency explicitly

Server Clusters


- ▶ In order to support Internet-wide services, we need to use **server clusters/server farms**.
- ▶ A simple approach is to route the requests at the TCP level:



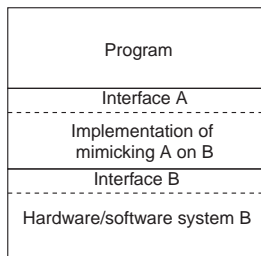
- ▶ The crux is to balance the load on the different servers
 - ▶ **Round-robin** is perhaps the simplest approach
 - ▶ Application-layer solutions are also possible

Resource Virtualization (1/3)

Idea



(a)




(b)

Essentially, virtualization allows to emulate the behavior of a different system.


Objective Allow legacy SW developed for a specific mainframe architecture to execute on a different platform (IBM)

Resource Virtualization (2/3)

Implementations



(a)

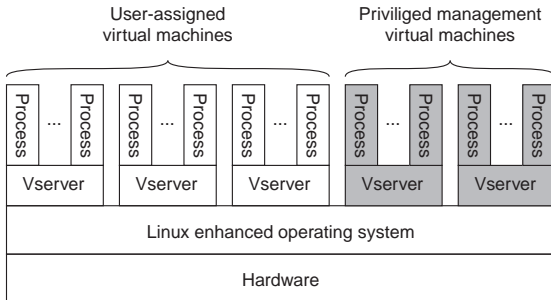


(b)

Process Virtual Machine (a) Each VM provides an instruction set and a run-time to support the execution of a single application. Example: *Java Virtual Machine*.

Virtual Machine Monitor (b) Each VM provides an instruction set and a run-time that allows the concurrent and independent execution of multiples OSs. Example: VMware, Xen, VirtualBox.

Resource Virtualization (3/3)



- ▶ The use of VMs in distributed systems has mainly two advantages:
 1. Improves security and reliability.
 2. Simplifies management

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Servers and State (1/4)

Problem the execution of the same task on every request may unnecessarily tax the server

Solution the server can keep some **state (information)**, i.e. information about the status ongoing interactions with clients;

- ▶ the size
- ▶ the processing demands

of each message are potentially smaller

- ▶ For example, in a distributed file system, the server may avoid open and close a file for each remote read/write operation
 - ▶ The server may keep a cache of open files for each client
- ▶ Depending on whether or not a server keeps state information, a server is called **stateful** or **stateless**, respectively

Servers and State (2/4)

- ▶ Keeping state information raises some challenges:
 - ▶ of consistency;
 - ▶ of resource management;upon failure of either clients or server
- ▶ Loss of state when a server crashes may lead to:
 - ▶ ignoring or rejecting client requests after recovery:
 - ▶ the client will have to start a new **session**
 - ▶ wrong interpretation of client requests sent before the crash:
 - ▶ TCP connection port reuse
- ▶ Keeping state (on server) when the client crashes may lead to:
 - ▶ resource depletion
 - ▶ wrong interpretation of requests sent by other clients after the crash

Servers and State (3/4)

- ▶ But not keeping state information in the server does not solve the problems arising from failures:
 - ▶ message duplication may lead to handling the same request several times
 - ▶ requests must be **idempotent**, if the transport protocol does not ensure non-duplication of packets;
 - ▶ the outcome may not be that satisfactory, if the transport protocol ensures non-duplication of packets

Servers and State(4/4)

- ▶ **Obs.-** Statelessness is a protocol issue:
 - ▶ A server can be stateless only if each protocol message has all the information for its processing independently of previous communication;
 - ▶ Inversely, a server can be stateful only if each protocol message has enough information to relate it to previous communication
- ▶ For example, Netscape had to add HTTP-header fields specifically for **cookies**.
 - ▶ HTTP is essentially stateless
 - ▶ Cookies are a device that allows a server to keep state about a client session:
 - ▶ *cookies* are stored on the client side

Client Identification in Stateful Servers

1. Use the address of the **access point**, i.e. of the channel endpoint
 - ▶ For example, the client's IP address and port
 - ▶ Issue: may not be valid for more than one transport session:
 - ▶ E.g. if a TCP connection breaks and a new one is setup in its place, the port number on the client's side may be different
2. Use a transport-layer independent **handle**. For example:
 - ▶ HTTP cookies

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Failures

Challenges:

1. components in a distributed application may fail, while others continue operating normally
2. on the Internet it is virtually impossible to distinguish network failures from host failures

Solution: highly application dependent

Client Crashes

Challenge resources reserved for the client may remain allocated forever

- ▶ sockets, for connection based communication
- ▶ state, in the case of stateful servers
- ▶ application specific resources

Solution **leases** (and timers):

- ▶ a server *leases a resource* to a client for only during a finite time interval: upon its expiration, the resource may be taken away, unless the client **renews** the lease

Server Crashes

Problem I: server may loose state

- ▶ may accept duplicated messages after recovery;

Problem II: how can we ensure that a request was performed?

- ▶ You are at an ATM. You type your PIN. You choose to withdraw 50 Euros. Suddenly, the machine reports communication error (or was it a server crash?) and does not give you the money. Was it withdrawn from your account?

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Security

Challenge: servers execute with privileges that their clients usually do not have

Solution: servers must

authenticate clients: i.e. "ensure" that a client is who it claims to be;

control access to resources: i.e. "ensure" that a client has the necessary permissions to execute the operation it requests.

- ▶ A related requirement is data **confidentiality**
 - ▶ need to encrypt data transmitted over the network
- ▶ Code migration (i.e. downloaded from the network) raises even more issues.

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Communication Channel Adaptation

Order the application will have to reorder the messages (must use a sequence number), if that is important

Reliability need to use timers to recover from message loss. Have to be aware of the possibility of duplicates.

Flow control: if you want to avoid message loss because of insufficient resources

Channel abstraction: the application may have to build messages from a stream. Or, fragment messages at one end and reassemble them at the other end.

Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

Further Reading

- ▶ Ch. 3 of Tanenbaum e van Steen, *Distributed Systems, 2nd Ed.*
 - ▶ Subsection 3.1.2 *Threads in Distributed Systems*, we assume the remaining material in Section 3.1 to be background knowledge (OS class)
 - ▶ Subsection 3.3.2 *Client-Side Software for Distribution Transparency*
 - ▶ Section 3.4 *Servers*
 - ▶ Section 3.2 *Virtualization*
- ▶ Arpaci-Dusseau & Arpaci-Dusseau, *Event-based Concurrency*, Ch. 33 of OSTEP book
- ▶ Pai et al., *Flash: An efficient and portable Web Server*, in 1999 Annual Usenix Technical Conference
- ▶ Welsh et al, *SEDA: An Architecture for Well-Conditioned, Scalable Internet Services*, in Symposium on Operating Systems, 2001